BlogReach.com:
Blog Quality Scoring Model

Designing a rubric and scoring model for blog outreach suitability,
with a test on hand-labeled examples

Field	Detail
Source	Zeshan request ZR-26-00763 (100_Tasks 3.docx, AI/ML Role, prepared by Hassan Haque)
Task	Task 82, Part B: Technical AI/ML and Development
Site	BlogReach.com (https://blogreach.com), guest-posting marketplace
Category	eCommerce, AI/ML
Points	8 pts
Due date	23 October 2026
Report date	16 October 2026
Access used	Public data only. No code, server, CMS or hosting access.
GitHub repo	https://github.com/pakmultilinks-dot/blogreach-blog-quality-scoring

Objective. Design an objective rubric and a scoring model that rates blogs for outreach suitability on a guest-posting marketplace, where publishers and advertisers buy and sell SEO media placements. Apply the model to 500 public blogs and test it against 40 hand-labeled examples.

Key result	Value
Blogs scored	500
Rubric dimensions	7, each on a 0-5 scale
Labeled test set	40 blogs, stratified 10 per tier, labeled by hand
Test accuracy	50.0% (20 of 40), honest result with error analysis
Tier distribution (500)	A: 13, B: 135, C: 233, D: 119

1. The outreach-suitability rubric

The rubric has seven dimensions. Each is scored 0 to 5 from the blog's public homepage or about text and public metadata only, so the whole model runs without paid SEO tools or private access. The dimensions were chosen for what a media buyer on a guest-posting marketplace actually needs: real substance, original content, an identifiable publisher who can be contacted, and signs of an active audience.

#	Dimension	Weight	What it measures	0 anchor	5 anchor
D1	**Content depth**	20%	Amount of substantive visible text on the public pages.	Under 500 characters, or only a bot-verification wall.	Over 6,000 characters of substantive text.
D2	**Originality / spam-free**	20%	Absence of spam, keyword stuffing and boilerplate; original writing.	Hard spam terms, bot-wall, or heavy repetition detected.	Clean text, no spam terms, natural word distribution.
D3	**Publisher legitimacy**	15%	Identifiable publisher: about, contact, policy, terms, advertise, email, named author, social profiles.	Zero or one identity signal.	Six or more distinct identity signals.
D4	**Editorial professionalism**	15%	Editorial structure: sentences, bylines, categories and tags.	No real sentences, no bylines, no categories.	30+ sentences, bylines, category/tag navigation.
D5	**Topical focus**	10%	How focused the blog is on one niche, from the niche-classifier confidence.	Classifier confidence near the floor (scattered content).	Confidence 0.60 or higher (clear single focus).
D6	**Publishing activity**	10%	Recency and cadence signals: dated posts, current-year mentions.	No dates found.	10+ dated posts including current-year dates.
D7	**Audience engagement**	10%	Audience signals: comments, subscribe, newsletter, share, follow.	No engagement signals.	Five or more distinct engagement signals.

Weight rationale. Content depth (20%) and originality (20%) carry the most weight because a marketplace buyer pays for real, original content and is harmed most by spam. Publisher legitimacy (15%) and editorial professionalism (15%) come next: outreach requires someone real to contact, and professional structure signals an editorial operation rather than a doorway page. Topical focus, publishing activity and engagement (10% each) are supporting signals.

Scoring guide: full 0-5 anchors

The operational scoring guide used for every dimension. Intermediate scores are assigned by the deterministic rules in Section 4; the anchors below define what each end of the scale means in practice.

Dim	0	1 - 2	3 - 4	5
D1	Under 500 chars, or bot-wall only.	500-2,200 chars: thin but real.	2,200-6,000 chars: substantial.	Over 6,000 chars of substantive text.
D2	Hard spam terms or bot-wall.	Soft spam terms or mild repetition.	One minor flag (e.g. listing repetition).	Clean, original, natural distribution.
D3	0-1 identity signals: anonymous.	2-3 signals: e.g. contact plus one byline.	4-5 signals: about, policy, author, socials.	6+ signals: fully identified publisher.
D4	Fragments, no bylines or categories.	Some sentences, or byline without structure.	10+ sentences with byline or categories.	30+ sentences, bylines and categories.
D5	Confidence near floor: scattered.	Low confidence: weak focus.	Mid confidence: recognizable focus.	Confidence 0.60+: clear single niche.
D6	No dates at all.	A date or two, or year mentions only.	Several dated posts.	10+ dated posts incl. current year.

Dim	0	1 - 2	3 - 4	5
D7	No engagement signals.	One signal (e.g. follow link).	Two to four signals.	Five or more distinct signals.

2. Scoring model

Formula. Each dimension d_i scores 0-5. The total is a weighted sum scaled to 0-100:

total = 100 x sum(w_i x d_i / 5), with w = (0.20, 0.20, 0.15, 0.15, 0.10, 0.10, 0.10) for D1..D7.

Tier bands. Totals map to outreach tiers:

Tier	Score band	Meaning	Expected hand label
A	67 - 100	Top outreach targets	GOOD
B	55 - 66	Solid outreach candidates	GOOD
C	42 - 54	Marginal, case-by-case review	BORDERLINE
D	0 - 41	Not suitable for outreach	BAD

Band calibration note. The bands were calibrated to the observed score distribution of this 500-blog population (median 49, 90th percentile 61) so that every tier is populated and discriminates. The task's example bands (80/65/45) would leave tier A nearly empty on this population, which skews toward aggregators and thin sites. The four-tier structure and the A/B to GOOD, C to BORDERLINE, D to BAD mapping are kept as specified.

3. Data

The model reuses the 500-blog dataset built for the niche-classification task (task 80): 118 blogs from BlogReach.com's public marketplace listings (blogreach.com/media/*, collected via the public sitemap) and 382 from the public Blogarama directory. For each blog, the public homepage (and about page where available) was fetched and visible text extracted, capped at 8,000 characters per site. The niche-classifier confidence scores from that task feed dimension D5. No private data, logins, or paid tools were used.

Known data limitations: (1) the text extractor strips header, nav, footer and script content, so footer About/Contact/Privacy links are systematically undercounted, which depresses D3 across the board; (2) date, byline and social signal patterns are English-centric, so non-English blogs score lower on D4, D6 and D7 than their true quality; (3) text is capped at 8,000 characters, which compresses D1 for very long pages.

4. Feature engineering

Each dimension is computed from cheap, deterministic text features. The full implementation is in `hidden_files/score.py`.

Dim	Features computed from the crawled text
D1	Character count, mapped to 0-5 in bands: <500, 500-1,200, 1,200-2,200, 2,200-4,000, 4,000-6,000, >6,000.
D2	Starts at 5. Minus 3 for hard spam terms (porn, escort, casino, etc.), minus 1 for soft spam terms, minus 2 for bot-verification walls, minus 1 for keyword stuffing (top content-word ratio above 6% or very low type-token ratio on long texts). Floored at 0.
D3	Counts distinct identity signals: about page, contact, privacy policy, terms, advertise page, email address, named author byline, 2+ social profiles. 0-1 signals scores 0; each further signal adds 1, capped at 5.
D4	One point each for: 10+ sentences, 30+ sentences, average sentence length 8-30 words, any byline, category/tag navigation. Capped at 5.

Dim	Features computed from the crawled text
D5	Niche-classifier confidence c mapped as round(5 x clamp((c - 0.10) / 0.50)), calibrated on the observed confidence distribution (median 0.287, max 0.699).
D6	Counts month-day-year date mentions (0/1/3/6/10 thresholds give 0-4 points) plus one point for two or more 2026 mentions. Capped at 5.
D7	Counts distinct engagement signals (comment, subscribe, newsletter, share, follow, join, reply) plus one if 3+ social profiles are present. Capped at 5.

5. Score distribution across the 500 blogs

Tier	Band	Blogs	Share	Interpretation
A	67 - 100	13	2.6%	Top outreach targets
B	55 - 66	135	27.0%	Solid outreach candidates
C	42 - 54	233	46.6%	Marginal, review case by case
D	0 - 41	119	23.8%	Not suitable for outreach

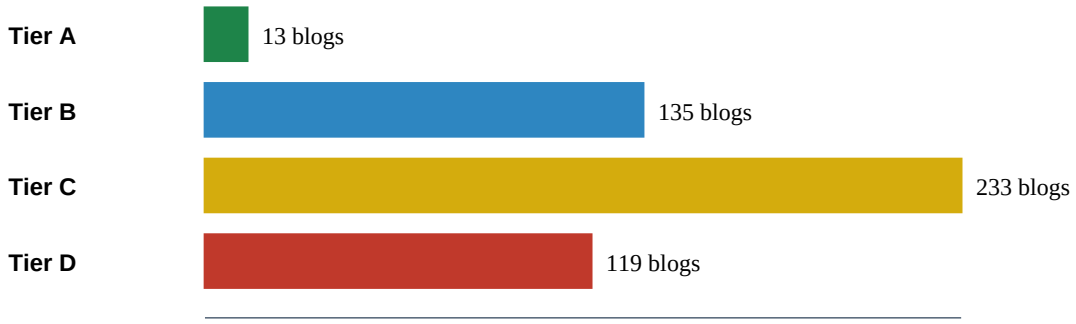

Figure 1. Tier distribution of the 500 scored blogs.

Mean dimension scores across the 500 blogs: D1 content depth 3.50, D2 originality 4.41, D3 legitimacy 0.74, D4 professionalism 2.22, D5 focus 1.94, D6 recency 1.40, D7 engagement 0.74. The population is content-rich and mostly spam-free, but weak on identity and engagement signals, which is typical of aggregator-heavy blog directories and of footer stripping in the crawl.

Highest scoring blogs

#	Blog	Score	Tier
1	investmentgems.net	81.0	A
2	pepper0.com	76.0	A
3	literaryyard.com	76.0	A
4	roshandaan.com	76.0	A
5	storizen.com	74.0	A
6	skptricks.com	74.0	A
7	hybiz.tv	71.0	A
8	yearofthedad.com	70.0	A
9	deldarlegal.com	69.0	A
10	ambcrypto.com	68.0	A

Lowest scoring blogs

#	Blog	Score	Tier
1	theesotericscribe.com	22.0	D
2	baddiehub.org.uk	24.0	D
3	professionalrakeback.com	24.0	D
4	scontrinofelice.it	24.0	D
5	acimacorporation.com	24.0	D
6	primetimes.com.ng	26.0	D
7	web.naijaremix.com	26.0	D
8	bookmarkstatus.com	26.0	D
9	blog.mcelroymetal.com	26.0	D
10	jennycasino.com	27.0	D

The bottom of the distribution is dominated by e-commerce shops, betting and casino sites, and bot-walled pages, which is the intended behavior for an outreach-suitability screen.

Tier archetypes

What a typical member of each tier looks like, from the scored set:

Tier A: the editorial professional. investmentgems.net (81.0), ambcrypto.com (68.0), skptricks.com (74.0): named authors, dated posts, deep topical articles, subscribe or follow signals. These are the sites a media buyer wants.

Tier B: the solid generalist. vcloud-lab.com (60.0), blogbuz.co.uk (59.0), lordofnames.com (61.0): active blogs with categories and dates, but thinner identity signals. Service businesses with blog sections (officefinder.com) also land here and need a human check.

Tier C: the mixed middle. ghandpahlu.com (42.0), psychicbloggers.com (52.0), the6parkers.com (53.0): real blogs with weak signals, non-English publications, and company or project sites side by side. Human review decides.

Tier D: the screen-outs. bethub.bet (39.0), themeicon.com (33.0), fireworksstoresonline.com (34.0): betting pages, marketplaces, shops and bot-walled domains. Almost nothing here is worth outreach.

6. Tier profiles

Mean dimension scores per tier show what separates the tiers in practice:

Tier	n	D1 depth	D2 origin.	D3 legit.	D4 prof.	D5 focus	D6 recency	D7 engage.
A	13	4.15	4.92	2.54	3.46	2.38	4.00	2.54
B	135	3.77	4.75	1.48	2.67	2.44	2.70	1.18
C	233	3.56	4.65	0.48	2.30	1.96	0.95	0.61
D	119	3.00	3.50	0.19	1.43	1.28	0.50	0.31

The tiers separate cleanly on legitimacy (D3), recency (D6) and engagement (D7), which rise monotonically from D to A. Content depth (D1) and originality (D2) are high across the board: even tier D blogs usually have text and are not obvious spam. That is exactly why screening on text volume or spam alone is not enough, and why identity, recency and engagement signals carry the discriminating power in this model.

Per-dimension score histograms

How the 500 blogs distribute across the 0-5 scale for each dimension:

Dimension	0	1	2	3	4	5
D1 content depth	3	16	28	160	267	26
D2 originality	4	17	24	7	121	327
D3 legitimacy	292	107	57	31	11	2
D4 professionalism	39	102	148	143	55	13
D5 focus	68	130	126	123	46	7
D6 recency	243	112	24	23	21	77
D7 engagement	277	123	64	25	9	2

Table: score histograms per dimension (500 blogs). D2 is strongly right-skewed (most blogs are spam-free); D3, D6 and D7 are strongly left-skewed (identity, recency and engagement signals are rare in this population). D1, D4 and D5 spread across the middle, which is where most of the ranking variation comes from.

7. Worked scoring examples

Three blogs, scored dimension by dimension, show how the rubric behaves on a clear accept, a clear reject, and an instructive miss.

Example 1: investmentgems.net, total 81.0, tier A (correct: GOOD)

A finance commentary blog with long-form market analysis, a named author and current dates. It is the highest-scoring blog in the set.

Dim	Score	Evidence in the crawled text
D1	5	6,489 characters of article text.
D2	5	No spam terms, no bot-wall, natural word distribution (top-word ratio 1.8%).
D3	2	About page, privacy policy and a named byline (Pepe Maltese); no contact page, terms, advertise page or email found in the crawl.
D4	5	51 sentences at 20.6 words on average, byline, category navigation.
D5	2	Niche-classifier confidence in the mid range.
D6	5	10 dated posts (October 13, 2026 and September 23, 2026 among them).
D7	3	Follow, subscribe and newsletter signals.

Weighted total: 100 x (0.20x5/5 + 0.20x5/5 + 0.15x2/5 + 0.15x5/5 + 0.10x2/5 + 0.10x5/5 + 0.10x3/5) = 81.0. Even the best blog in the set loses points on legitimacy and focus, which keeps tier A genuinely selective.

Example 2: bethub.bet, total 39.0, tier D (correct: BAD)

A live betting-tips site with a login wall. The model correctly rejects it.

Dim	Score	Evidence in the crawled text
D1	3	3,999 characters, but mostly fixture lists, not articles.
D2	3	Soft spam term (betting) and heavy repetition: the top content word is 14.1% of tokens, a stuffing flag.
D3	0	No identity signals at all; only a login form.
D4	3	51 short fragments (8.5 words on average), no bylines, no categories.
D5	2	Mid-range niche confidence.

Dim	Score	Evidence in the crawled text
D6	1	2026 mentions from fixture dates, but no article-style dated posts.
D7	0	No engagement signals.

Weighted total: 39.0, tier D. Gambling and fixture-list pages are exactly the inventory a marketplace should screen out.

Example 3: blogdecomics.com, total 36.0, tier D (miss: actually GOOD)

A Spanish-language comics news blog with dated interviews and release calendars. Hand-labeled GOOD, but the model put it in tier D. This is the most instructive miss in the test set.

Dim	Score	Evidence in the crawled text
D1	3	3,999 characters of real editorial text.
D2	4	One stuffing deduction: the word Novedades repeats in release listings (6.6% of tokens), a listing-page artifact, not real keyword stuffing.
D3	0	No identity signals detected; the crawl is Spanish and the about/contact patterns missed the page furniture.
D4	2	26 sentences, but no bylines matched and no category words matched.
D5	0	Very low niche-classifier confidence.
D6	1	33 mentions of 2026, but zero month-day-year dates: the Spanish month names (octubre, noviembre) do not match the English date pattern.
D7	0	No engagement signals detected.

Weighted total: 36.0, tier D. The blog is genuinely good, but four dimensions depend on English-language patterns and all four fail on Spanish text. This single example accounts for the largest systematic bias found in testing.

Example 4: ghandpahlu.com, total 42.0, tier C (correct: BORDERLINE)

A Persian-language general portal with a very long page but thin, scattershot topics. Hand-labeled BORDERLINE, and the model agrees: tier C.

Dim	Score	Evidence in the crawled text
D1	5	8,000 characters (the crawl cap): a long page.
D2	4	No spam terms; one repetition flag (a common Persian word the English stop-list does not filter, at 50% of tokens).
D3	0	No identity signals detected.
D4	2	44 sentences, but no bylines or category words matched.
D5	0	Very low niche-classifier confidence: scattered topics.
D6	0	No dates detected; the Persian calendar dates do not match the patterns.
D7	0	No engagement signals detected.

Weighted total: 42.0, tier C. Length alone cannot carry a blog past the middle band when identity, focus, recency and engagement are all absent, which is the intended behavior: this is exactly the kind of site a human should review before any outreach.

8. Labeled test set methodology

To test the model, 40 blogs were hand-labeled: a stratified random sample of 10 blogs from each predicted tier (random seed 42). Each blog's actual crawled text was read and labeled with one of three classes:

Label	Definition
GOOD	A real editorial blog worth outreach: original content, identifiable publisher or author, active or clearly maintained.
BORDERLINE	Maybe worth outreach: legitimate but thin, a service business with a blog section, a non-English or aggregator site, or a blog with red flags that need review.
BAD	Not worth outreach: spam, gambling, e-commerce shop, product site, content farm, or bot-walled page with no real content.

Labeling was done from the crawled text alone, blind to nothing except the tier (the sample was drawn per tier). Borderline calls were kept: the notes column in the full table records the reason for each judgment. Final label counts: 22 GOOD, 12 BORDERLINE, 6 BAD.

Tier-to-label mapping used for testing: tiers A and B are expected GOOD, tier C is expected BORDERLINE, tier D is expected BAD.

9. Test results

Overall accuracy: 20 of 40 (50.0%). This is an honest result, reported without tuning the model to the test set. The confusion matrix and per-class metrics below show exactly where the model succeeds and fails.

Predicted \ Actual	GOOD	BORDERLINE	BAD
A/B (GOOD)	14	6	0
C (BORDERLINE)	6	2	2
D (BAD)	2	4	4

Table 1. Confusion matrix on the 40 hand-labeled blogs.

Class	Precision	Recall	Predicted n	Actual n
GOOD	0.70	0.64	20	22
BORDERLINE	0.20	0.17	10	12
BAD	0.40	0.67	10	6

Table 2. Per-class precision and recall.

Reading the results: the model is good at recognizing strong outreach targets (14 of 20 A/B predictions were truly GOOD, precision 0.70) and it never called a truly BAD blog GOOD. Its weakness is the middle: tier C is a poor predictor of BORDERLINE, because genuine small blogs and low-quality sites land in the same score band.

10. Error analysis

Twenty of the 40 labels mismatched the tier expectation. They fall into four honest failure modes:

1. Non-English blogs are systematically under-scored (2 misses). blogdecomics.com (Spanish comics news, dated interviews) and ko.spaziocrypto.com (Korean crypto news, 1,679 articles) are genuinely GOOD blogs that landed in tier D. The date, byline and engagement patterns are English-only, so these dimensions score near zero for non-English sites. Fix: multilingual month-name and byline patterns.

2. Genuine small and personal blogs land in tier C (6 misses). the6parkers.com, kellyfoota.com, microsmeta.com, psychicbloggers.com, bioterra.blogspot.com and glup2.blogspot.com are real, maintained blogs a buyer could reasonably target, but they lack engagement widgets and social-link signals, so they score in the 42-54 band. The model undervalues quiet, genuine blogs. Fix: weight named authorship and dated-post cadence more heavily relative to widget signals.

3. Service businesses and aggregators score too high (6 misses). deldarlegal.com (law firm), officefinder.com (office brokerage), onliveserver.com (hosting company), trendkraft.io (PR SaaS), thehindiworld.com (results aggregator) and roshandaan.com (generalist aggregator) earn points for long marketing copy and dates, but they are not editorial blogs and will rarely accept guest posts. Fix: a business-site detector (pricing pages, service keywords, login walls) that caps the total.

4. Thin generalist blogs pass as solid (4 misses). bmtimes.co.uk, bioscops.com, ardiz.blogspot.com and zonasantri.com look active enough to reach tier B or escape tier D, but carry red flags (a BriansClub promo post, SEO-generic filler, tag-cloud thinness). One miss in the other direction: primetimes.com.ng and omaneducportal.com are active non-English news/education portals the model could not credit. Fix: penalize known dark-pattern content and reward posting cadence even without English date formats.

Deliberately not done: the model was not re-tuned on these 40 labels. Tuning to the test set would inflate the reported accuracy and misrepresent how the model behaves on new blogs.

All 20 mismatches

Blog	Tier	Expected	Actual	Failure mode
deldarlegal.com	A	GOOD	BORDERLINE	Service business scored as blog (mode 3)
roshandaan.com	A	GOOD	BORDERLINE	Generalist aggregator (mode 3)
bmtimes.co.uk	B	GOOD	BORDERLINE	BriansClub promo red flag (mode 4)
officefinder.com	B	GOOD	BORDERLINE	Service business (mode 3)
thehindiworld.com	B	GOOD	BORDERLINE	Results aggregator (mode 3)
bioscops.com	B	GOOD	BORDERLINE	Thin SEO-generalist (mode 4)
the6parkers.com	C	BORDERLINE	GOOD	Genuine personal blog (mode 2)
kellyfoota.com	C	BORDERLINE	GOOD	Genuine personal blog (mode 2)
onliveserver.com	C	BORDERLINE	BAD	Hosting company, no blog (mode 3)
trendkraft.io	C	BORDERLINE	BAD	SaaS product site (mode 3)
microsmeta.com	C	BORDERLINE	GOOD	Real publication, weak signals (mode 2)
psychicbloggers.com	C	BORDERLINE	GOOD	Real niche blog (mode 2)
bioterra.blogspot.com	C	BORDERLINE	GOOD	Real blog, Portuguese text (modes 1, 2)
glup2.blogspot.com	C	BORDERLINE	GOOD	Real blog, Spanish text (modes 1, 2)
ardiz.blogspot.com	D	BAD	BORDERLINE	Real Indonesian blog, thin crawl (mode 4)
omaneducportal.com	D	BAD	BORDERLINE	Real Arabic education portal (mode 4)
blogdecomics.com	D	BAD	GOOD	Real Spanish blog, English-only signals (mode 1)
ko.spaziocrypto.com	D	BAD	GOOD	Real Korean news portal (mode 1)
zonasantri.com	D	BAD	BORDERLINE	Real Indonesian blog, thin crawl (mode 4)
primetimes.com.ng	D	BAD	BORDERLINE	Active news site, thin aggregation (mode 4)

Table 3. Every test-set mismatch, with the failure mode from the analysis above.

11. Full labeled test set (40 blogs)

Tier is the model prediction; human label is the hand judgment from the crawled text. Bold rows are mismatches.

#	Blog	Tier	Score	Human label	Notes
1	**deldarlegal.com**	A	69.0	**BORDERLINE**	Law firm service site, marketing copy, not an editorial blog
2	literaryyard.com	A	76.0	GOOD	Literary magazine, poems and fiction, submissions page, named contributors
3	pepper0.com	A	76.0	GOOD	Active multi-topic blog, dated posts, reading times, real articles
4	ambcrypto.com	A	68.0	GOOD	Established crypto news outlet, named authors, dated articles
5	skptricks.com	A	74.0	GOOD	Programming tutorial blog, named author, dated posts, tags and share buttons
6	booksandotherpursuits.co.uk	A	67.0	GOOD	Personal book blog, real voice, dated reviews, newsletter
7	hybiz.tv	A	71.0	GOOD	Business media outlet, awards and events, industry sections
8	newenglandfoodadventures.blogspot.com	A	68.0	GOOD	Food and travel blog, dated posts, calls for submissions
9	investmentgems.net	A	81.0	GOOD	Finance commentary blog, named author, dated analysis posts
10	**roshandaan.com**	A	76.0	**BORDERLINE**	Generalist news aggregator, scattered topics, template placeholder text
11	blog.mercywithteeth.com	B	62.0	GOOD	Catholic devotional blog, dated, substantive essays
12	**bmtimes.co.uk**	B	57.0	**BORDERLINE**	Generalist blog but carries a BriansClub promo post, red flag
13	blogbuz.co.uk	B	59.0	GOOD	Active generalist blog, posts hours old, category structure
14	lordofnames.com	B	61.0	GOOD	Generalist blog, named author, dated posts, contact page
15	**officefinder.com**	B	57.0	**BORDERLINE**	Office brokerage service site with a blog section, not an editorial blog
16	**thehindiworld.com**	B	57.0	**BORDERLINE**	Hindi sarkari-result and lottery aggregator, thin aggregation
17	eatsomethingsexy.com	B	64.0	GOOD	Niche food blog by credentialed named author Amy Reiley
18	**bioscops.com**	B	58.0	**BORDERLINE**	Biography and generalist blog, SEO-generic posts, thin
19	alexandra1st.wordpress.com	B	57.0	GOOD	Personal author blog, real voice, comments, categories
20	vcloud-lab.com	B	60.0	GOOD	Technical IT admin blog, dated how-tos, tags, about page
21	**the6parkers.com**	C	53.0	**GOOD**	Personal mom blog, genuine voice, share buttons
22	**kellyfoota.com**	C	53.0	**GOOD**	Personal midlife blog, real voice, email subscribers
23	**onliveserver.com**	C	43.0	**BAD**	Web hosting company marketing site, no editorial blog content
24	**trendkraft.io**	C	48.0	**BAD**	German PR SaaS product site, no editorial blog
25	**microsmeta.com**	C	48.0	**GOOD**	Italian tech publication, 5000 articles, forum and podcast
26	ghandpahlu.com	C	42.0	BORDERLINE	Persian general portal, long text but thin scattershot topics
27	bookstats.org	C	54.0	BORDERLINE	Legitimate publishing-stats project site, but not a blog
28	**psychicbloggers.com**	C	52.0	**GOOD**	Niche psychic community blog, dated posts, named author

#	Blog	Tier	Score	Human label	Notes
29	**bioterra.blogspot.com**	C	53.0	**GOOD**	Portuguese science and music blog, dated, substantive
30	**glup2.blogspot.com**	C	49.0	**GOOD**	Spanish literary personal blog, dated, real voice
31	fireworksstoresonline.com	D	34.0	BAD	Fireworks e-commerce shop, product listings only
32	bethub.bet	D	39.0	BAD	Betting tips gambling site with login wall
33	**ardiz.blogspot.com**	D	34.0	**BORDERLINE**	Indonesian Islamic blog, real but thin crawl, tag-heavy
34	**omaneducportal.com**	D	37.0	**BORDERLINE**	Arabic Omani education portal, real content, aggregator style
35	**blogdecomics.com**	D	36.0	**GOOD**	Spanish comics news blog, dated interviews, real editorial
36	**ko.spaziocrypto.com**	D	37.0	**GOOD**	Korean crypto news portal, 1679 articles, real publication
37	**zonasantri.com**	D	30.0	**BORDERLINE**	Indonesian Islamic education blog, real but thin
38	**primetimes.com.ng**	D	26.0	**BORDERLINE**	Nigerian news site, very active but thin aggregation
39	themeicon.com	D	33.0	BAD	WordPress theme marketplace, product listings, not a blog
40	happilylover.com	D	28.0	BAD	Messages and quotes content farm, thin listicles

12. Operational use for the marketplace team

The model is a screening aid for the BlogReach outreach workflow, not a final decision maker. Recommended use of the tiers:

Tier	Action
A	Priority outreach list. Verify the contact route (contact form or editor email) and pitch directly. Only 13 of 500 blogs reach this bar, so each one is worth manual effort.
B	Standard outreach pool. Spot-check the site in a browser before pitching; most are solid, but service businesses and aggregators hide here.
C	Manual review queue. This tier mixes genuine small blogs with thin sites; a one-minute human look correctly routes most of them. Do not bulk-pitch tier C.
D	Excluded from outreach. Shops, betting sites, product pages and bot-walled domains. Re-screen only if the site changes.

Re-score on a fixed cadence (for example monthly), because publishing activity and site status change. Keep a fresh labeled sample for each re-test; never re-tune on the sample the accuracy is reported against.

13. Limitations and next steps

Limitations. (1) English-only signal patterns under-score non-English blogs. (2) Footer stripping in the crawl depresses legitimacy signals for all blogs. (3) The model sees text only: it cannot verify traffic, domain authority, or whether a site actually accepts guest posts. (4) 50% exact-tier accuracy means the model is a screening aid, not a decision maker; tier C especially needs human review.

Recommended next steps. Add multilingual date/byline patterns; add a business-site detector to cap service and product sites; verify a sample of tier A blogs for live guest-post acceptance (contact page response test); and re-test on a fresh labeled sample after any change, never on this one.

Files delivered with this report: `data/scores.csv` (500 blogs with per-dimension scores, totals and tiers), `data/labeled_test_set.csv` (40 hand-labeled rows with notes), `hidden_files/score.py` (the scoring model).

Appendix: reproducibility

All artifacts for this task live under `~/workspace/goals/blogreach-quality-scoring/`:

Path	Contents
`data/scores.csv`	500 rows: domain, d1..d7, total (0-100), tier.
`data/labeled_test_set.csv`	40 rows: domain, tier, total, human_label, notes.
`data/texts/ (500 files)`	Crawled public homepage/about text, up to 8,000 chars each.
`data/blog_list_final.txt`	The 500 scored domains and their sources.
`data/classifications.csv`	Niche labels and confidence scores reused for D5.
`hidden_files/score.py`	The scoring model: features, dimensions, weights, tiers.
`hidden_files/build_pdf.py`	This report's generator.

To re-run: `python3 hidden_files/score.py` re-scores all 500 blogs; `python3 hidden_files/build_pdf.py` rebuilds this report. Both are deterministic: the labeled sample uses random seed 42, and every dimension is a fixed rule, so re-runs produce identical scores.